

Practical RESTful API design

Exploration over technologies enabling a Model-Driven approach in
RESTful API development

Giovanni Ciatto

`giovanni.ciatto@studio.unibo.it`

University of Bologna
Computer Science and Engineering

May 18, 2016

Table of Contents

1 REST in practice

- Overview
- Principles
 - Client-server
 - Stateless interaction
 - Cacheability
 - Uniform Interface

1 REST in practice

- Overview
- Principles
 - Client-server
 - Stateless interaction
 - Cacheability
 - Uniform Interface

REST = REpresentational State Transfert

- REST *is* basically a concept, a set of principles (best practices) referred as “constraints” and described through natural language
- REST *is not* a tight specification, and there exists more than a way to produce RESTful web services
 - Yeah, this is about web-services, so cool!
- It's not a dichotomy, nor simply a matter of RESTful and RESTless web-services:
 - there are *more RESTful* ones ...
 - ... and *more RESTless* ones
- RESTful (distributed) systems tend to be scalable, robust and easy to develop, understand and maintain
 - Yet not every system should be RESTful
 - Silver-bullets are like unicorns: they do not exist.

REST constraints

RESTful systems **should** satisfy the following 6 constraints:

- Client-server
- Stateless interaction
- Chaceability
- Uniform Interface
- Layered System
- Code on Demand (Optional)

1 REST in practice

- Overview
- Principles
 - Client-server
 - Stateless interaction
 - Cacheability
 - Uniform Interface

Client-server I

There are two kind of entities: *servers* and *clients* communicating via a *connector*.

Servers

Reactive entities providing one or more services to multiple clients

Clients

Triggering (proactive) entities making requests that trigger reactions from servers.

Connectors

Mechanism that allows communication between clients and servers (e.g. HTTP protocol, Message Oriented Middleware, RPC, etc.)

Separation of concerns is the key concept here.

Separation of concerns

Once both clients and servers concerns have been fixed and some sort of common interface have been defined, the two kind of components evolve independently.

Client-server constraint in web-based systems

Web-based systems quite often satisfy this constraint.

Stateless interaction I

This is probably the most important constraint as well as the hardest one: each request from client to server must contain all of the information necessary to understand the request, and cannot take advantage of any stored context on the server. Application state is therefore kept entirely on the client.

Application state

Data that could vary by client, and per request.

Exceptions are tolerated

Immutability is a utopia. For real world problems, you should just try to minimize mutability. E.g. request rate monitoring requires mutability. However, make every effort to ensure that application state doesn't span multiple requests of your services.

Stateless interaction II

Cookies

Cookies do not necessarily violate this constraint, since they are part of each HTTP interaction.

Sessions identifiers

Assigning any sort of temporary session identifier to the some user *and* storing session data using some data structure within server's memory is inherently a violation of the constraint. Since this is the common usage of cookies, they are discouraged.

Example

A stateless communication litmus test is to turn off session cookies, and determine if the API, web service, or web application still works as designed.

Stateless interaction III

Authentication

This constraint makes authentication critical. Other approaches have been developed which are more secure and RESTful.

Cacheable responses

- Clients can cache responses
- Responses must therefore, implicitly or explicitly, define themselves as cacheable, or not
 - to prevent clients reusing stale or inappropriate data in response to further requests
- Well-managed caching partially or completely eliminates some client-server interactions
 - further improving scalability and performance

Uniform Interface I

- The uniform interface constraint defines the interface between clients and servers.
- This is probably secret of REST's simplicity and strength.
- RESTful systems expose a standard, unambiguous, clear and human-readable API because of this constraint.
 - Such a principle allows for model-driven approaches.

Uniform Interface Requirements

- Resource-based
- Manipulation of Resources Through Representations
- Self-descriptive Messages
- HATEOAS (dafuq ?!)

Resource-Based

- RESTful systems handle *resources*: servers host resources and clients want to CRUD them
 - ⇒ **C**reate or **R**ead or **U**ppdate or **D**eleate
- Resources have a *hierarchical* nature
- Resources are identified and referenced by the mean of URIs
 - ⇒ **U**niform **R**esource **I**dentifier

HTTP Verbs

- CRUD in Web-based systems means exploiting HTTP methods, which are often called “verbs” because of their usage within APIs:
 - POST is used for resource Creation
 - GET mean is to Read resources
 - PUT aim is to Update resources
 - DELETE is used to Delete resources
- HTTP Status codes and their general purpose semantics are part of the uniform interface too:
 - e.g. any successful request will result in a 200: Ok or 204: No Content status code (depending on whether the response has body or not)
 - e.g. trying to GET or DELETE any non-existent resource will result in a 404: Not Found status code
 - e.g. POSTing an already-existing user will result in a 409: Conflict status code

- Example of resource creation:
 - we want to edit some user's username from gciatto92 to gciatto
 - suppose no authentication is needed

Example of non-RESTful approach

- GET `http://example.com/users?user=gciatto92&-operation=changeName&newName=gciatto`
- ✗ RPC style: the request contains the to-be-called operation
- ✗ No semantics for GET verb
- ✗ Which resource am I editing?

Example of RESTful approach

- PUT `http://example.com/users/gciato92?newName=gciato`
- ✓ I'm editing the resource `gciato92`, which is a user, composing the `users` resource, which represents the collection of registered users
- ✓ PUT verb means Update and that's what I am doing
- ✓ URI queries are simply a mean for payload transport

Manipulation of Resources Through Representations

- Resources are not accessed (CRUDed) directly but through their representation(s)
- Representations should expose resources traits to clients enabling them to do what they are allowed to, no more and no less
- Clients cannot make assumptions upon resources implementations, they can only exploit representations

Self-descriptive Messages

- Each message must include enough information to describe how to process the message itself
 - E.g. HTTP's *Content Negotiation* is a powerfull feature: use it! Accept and Content-Type headers allow different representations for resources exposing some business logic
 - Prefer JSON but try to support XML

HATEOAS - **H**ypermedia **A**s **T**he **E**ngine **O**f **A**pplication **S**tate

- Difficult constraint to fully accomplish IMHO
- Weak definition is enough for our concerns: *“Services responses should contain ‘relational links’ when they may be useful”*
- Relationships are standardized
- In completely-RESTful HTTP-based systems, relational links are the only mean used by clients to interact with servers

Uniform Interface VIII

Example

GET <http://example.com/users?limit=3>

supposing the first three users are gciatto, mfrancia and mneri, it should return something like:

```
[
  {
    "username" = "gciatto",
    "links" = [
      { "rel" = "next", "href" = "/users/mfrancia" },
      { "rel" = "prev", "href" = "/users/mneri" },
      { "rel" = "first", "href" = "/users/gciatto" },
      { "rel" = "last", "href" = "/users/mneri" },
      { "rel" = "self", "href" = "/users/gciatto" }
    ]
  }
  ...
]
```



Fielding, Roy Thomas

Architectural Styles and the Design of Network-based Software Architectures.

Doctoral dissertation, University of California, Irvine, 200.